

# The Lottery Ticket Hypothesis

## Finding Sparse, Trainable Neural Networks

Aditya Mehta    Nikhil Goyal    Shardul Junagade

ES667: Deep Learning  
Indian Institute of Technology Gandhinagar

December 4, 2025

# Problem Description

# The Pruning Paradox

## Challenges in large networks:

- Overparameterized networks: millions to billions of parameters
- Deployment constraints: mobile devices, edge computing
- Post-training pruning works, but sparse networks fail when trained from scratch
- Difficult to find such optimal parameter to be pruned

### Key Question

If networks can be pruned by 90%+ after training, why can't we train smaller networks from scratch?

# The Lottery Ticket Hypothesis

Frankle & Carbin (2019)

*"A randomly-initialized dense network contains a subnetwork that is initialized such that, when trained in isolation, it can match the test accuracy of the original network."*

## Key Insights:

- Dense networks contain sparse subnetworks ("winning tickets")
- Same structure + random initialization = poor performance
- **Initialization matters!**

# Research Objectives

1. Validate winning ticket existence across architectures and datasets
2. Prove reinitialization to the original /starting weights is mandatory
3. Characterize sparsity ranges where LTH holds
4. Analyze learning dynamics and convergence while training
5. Examine the effect of various pruning techniques on generalization error

# Solution Approach

# Experimental Setup

## Datasets:

- MNIST (handwritten digits)
- Fashion-MNIST (clothing items)
- CIFAR-10 (natural images)

## Architectures:

- LeNet-300-100 (266K params)
- LeNet-5 (61K params)
- Conv-6 (1.2M params)

## Training Config:

- Optimizer: Adam
- Learning rate: 0.0012
- Iterations: 40,000
- Batch Size: 60
- Pruning rate: 30% \*
- Rounds: 16
- Pruning: IMP or Random
- Initialization: Rewind or Random

\* 30% of existing weights per layer

---

**Algorithm 1** Iterative Magnitude Pruning (IMP)

- 1: Initialize weights  $\theta_0 \sim \mathcal{D}$ ; store  $\theta_{init} \leftarrow \theta_0$ ; mask  $m \leftarrow \mathbf{1}$
  - 2: **for** round  $r = 1$  to  $N$  **do**
  - 3:   Train sparse network: optimize  $\theta$  using  $f(x; m \odot \theta)$  for  $T$  steps
  - 4:   **for** each layer  $\ell$  **do**
  - 5:     Compute pruning threshold  $\tau_\ell = p$ -th percentile of  $|\theta^{(\ell)}|$  among currently active weights
  - 6:     Update mask:  $m^{(\ell)} \leftarrow m^{(\ell)} \odot \mathbb{I}\{|\theta^{(\ell)}| > \tau_\ell\}$  (cumulative)
  - 7:   **end for**
  - 8:   Rewind:  $\theta \leftarrow m \odot \theta_{init}$
  - 9: **end for**
  - 10: Return sparse subnetwork  $(m, \theta_{init})$
- 

**Key:** Reset surviving weights to *original* initialization, not trained values!

## Three Experiment strategies:

1. **Magnitude + Rewind:** Prune low-magnitude weights, reset to  $\theta_0$

*Purpose: Find winning tickets*

2. **Magnitude + Random:** Same masks, reset to new random  $\theta'_0$

*Purpose: Test if initialization matters*

3. **Random + Rewind:** Random pruning, reset to  $\theta_0$

*Purpose: Test if structure alone works*

# Data Split and Evaluation Strategy

## Data Split:



- Training set:  $\sim 78.5\%$
- Validation set:  $\sim 7.1\%$
- Test set:  $\sim 14.3\%$

## Evaluation Procedure:

- Train the network across full epochs.
- After each epoch, compute validation loss.
- **Early Stop Iteration:** the iteration corresponding to the *minimum* validation loss.
- For each pruning round (40,000 iterations  $\approx 44$  epochs), test accuracy is reported at the early-stopped iteration.

# Results

# Winning Tickets Exist!



**Figure:** MNIST LeNet-5 (left) peaks at 99.16% with 49% sparsity. CIFAR-10 Conv-6 (right) improves 3.21% at 24% remaining weights!

# Winning Ticket Performance Summary

Dataset-Model	Peak Acc	Sparcity	Range*
MNIST-FC	98.35%	24%	24-6%
MNIST-Conv	<b>99.16%</b>	49%	49-6%
Fashion-FC	89.10%	17%	17-6%
Fashion-Conv	88.92%	49%	49-9%
CIFAR10-Conv6	86.30%	24%	24-6%

\*Key Finding: Networks maintain accuracy at 5% of original size!

# Initialization is Critical

Dataset-Model	Sparsity	Rewind	Random	Gap
MNIST-FC	1.98%	96.30%	76.55%	<b>19.75%</b>
MNIST-Conv	1.99%	94.31%	60.72%	<b>33.59%</b>
Fashion-FC	4.05%	86.25%	63.02%	<b>23.23%</b>

## Critical Finding

At extreme sparsity, random reinitialization causes **19-33% accuracy drops**. Structure alone is insufficient, initialization is essential!



# Magnitude Beats Random Pruning

Dataset-Model	Sparsity	Magnitude	Random	Gap
MNIST-FC	5.77%	97.79%	96.05%	1.74%
Fashion-Conv	5.78%	88.37%	84.27%	<b>4.10%</b>
CIFAR10-Conv6	5.77%	80.44%	77.15%	3.29%

## Critical Finding:

Magnitude-based weight selection is superior to random masking (1.74-4.10% advantage). Reinforces the idea that initialization rule matters more than pruning rule

# Learning Dynamics



## Observations:

- As the % of remaining parameters decreases, the model takes more iterations to converge
- Random reinitialization needs much more iterations as compared to rewind
- The trend is easily observable in simple datasets like MNIST, but much more noisy in case of complex datasets like CIFAR-10

# Conclusions

# Limitations

- LTH validated only on small vision **datasets**.
- LTH validated only using small models.
- Unstructured sparsity is not **hardware-friendly**.
- We only test for magnitude pruning and random pruning.

## Future Work

1. Apply LTH to **deeper architectures** (ResNets, VGG)
2. Test hypotheses for larger datasets
3. Structured pruning (**hardware-efficient pruning**)
4. Explore layer-wise adaptive pruning rates
5. **Early prediction** of winning tickets (before full training)

# Key Findings

1. **LTH validated:** Winning tickets exist across all architectures
  - Networks maintain 98% accuracy at 5-10% sparsity (MNIST)
  - Pruning can *improve* accuracy (3.21% on CIFAR-10)
2. **Regularizing effect:** We observe an implicit regularizing effect
3. **Initialization critical:** 19-33% drops with random reinit
4. **Magnitude pruning wins:** 1.74-4.10% better than random
5. **Dataset dependency:** LTH stability varies with task complexity

## Why Dense Networks Work

They contain multiple lottery tickets with favorable initializations!

# Thank You!

Questions?

GitHub: <https://github.com/aditya-me13/lottery-ticket-hypothesis>

# Appendix

# Detailed Dataset Specifications

Dataset	Details
<b>MNIST</b>	28×28 grayscale images, 10 classes (digits 0-9) Split: 55,000 train / 5,000 val / 10,000 test Normalization: $\mu = 0.1307$ , $\sigma = 0.3081$
<b>Fashion-MNIST</b>	28×28 grayscale images, 10 classes (clothing items) Split: 55,000 train / 5,000 val / 10,000 test Normalization: $\mu = 0.2860$ , $\sigma = 0.3530$
<b>CIFAR-10</b>	32×32 RGB images, 10 classes (natural objects) Split: 45,000 train / 5,000 val / 10,000 test Per-channel normalization: $\mu = (0.4914, 0.4822, 0.4465)$ $\sigma = (0.2023, 0.1994, 0.2010)$

# Model Architecture Details

## LeNet-300-100 (Fully Connected):

- Architecture: Input (784) → FC1 (300) → ReLU → FC2 (100) → ReLU → Output (10)
- Total parameters: ~266,000
- Used for: MNIST, Fashion-MNIST

## LeNet-5 (Convolutional):

- Conv1:  $1 \rightarrow 6 (5 \times 5)$  → ReLU → MaxPool ( $2 \times 2$ )
- Conv2:  $6 \rightarrow 16 (5 \times 5)$  → ReLU → MaxPool ( $2 \times 2$ )
- FC:  $256 \rightarrow 120 \rightarrow 84 \rightarrow 10$
- Total parameters: ~61,000
- Used for: MNIST, Fashion-MNIST

# Conv-6 Architecture

## Deep VGG-style Convolutional Network:

### Block 1:

- Conv1:  $3 \rightarrow 64$  ( $3 \times 3$ , padding=1)  $\rightarrow$  ReLU
- Conv2:  $64 \rightarrow 64$  ( $3 \times 3$ , padding=1)  $\rightarrow$  ReLU
- MaxPool ( $2 \times 2$ )

### Block 2:

- Conv3:  $64 \rightarrow 128$  ( $3 \times 3$ , padding=1)  $\rightarrow$  ReLU
- Conv4:  $128 \rightarrow 128$  ( $3 \times 3$ , padding=1)  $\rightarrow$  ReLU
- MaxPool ( $2 \times 2$ )

### Block 3:

- Conv5:  $128 \rightarrow 256$  ( $3 \times 3$ , padding=1)  $\rightarrow$  ReLU
- Conv6:  $256 \rightarrow 256$  ( $3 \times 3$ , padding=1)  $\rightarrow$  ReLU
- MaxPool ( $2 \times 2$ )

**Classifier:** FC:  $4096 \rightarrow 256 \rightarrow 10$

# Complete Training Configuration

## Optimization:

- Optimizer: Adam
- Learning rate: 0.0012
- Loss: Cross-entropy
- Batch size: 60 (128 for Conv-6)
- Total iterations: 40,000

## Pruning:

- Pruning rate: 30% per round
- Total rounds: 16
- Scope: Layer-wise
- Strategy: Magnitude-based
- Mask enforcement: After every gradient step

## Initialization:

- Method: Kaiming Normal
- Mode: fan\_in
- Non-linearity: ReLU

## Hardware:

- Platform: Kaggle Notebooks
- GPU: NVIDIA Tesla T4
- Framework: PyTorch 2.0+

- **Test Accuracy:** Final accuracy on test set after each pruning round
- **Validation Accuracy:** Used for early stopping criterion
- **Early-Stop Iteration:** Iteration number at minimum validation loss (proxy for convergence speed)
- **Sparsity Level:** Percentage of pruned weights

$$\text{Sparsity} = \left(1 - \frac{\|m\|_0}{\|\theta\|}\right) \times 100\%$$

- **Winning Ticket Range:** Pruning rounds where test accuracy remains within 2% of dense baseline
- **Accuracy Gap:** Difference between rewind and random reinitialization strategies